

# classification\_nn

April 25, 2024

## 1 ECE 285 Assignment 1: Classification using Neural Network

Now that you have developed and tested your model on the toy dataset set. It's time to get down and get dirty with a standard dataset such as cifar10. At this point, you will be using the provided training data to tune the hyper-parameters of your network such that it works with cifar10 for the task of multi-class classification.

Important: Recall that now we have non-linear decision boundaries, thus we do not need to do one vs all classification. We learn a single non-linear decision boundary instead. Our non-linear boundaries (thanks to relu non-linearity) will take care of differentiating between all the classes

TO SUBMIT: PDF of this notebook with all the required outputs and answers.

```
[1]: # Prepare Packages
import numpy as np
import matplotlib.pyplot as plt

from ece285.utils.data_processing import get_cifar10_data
from ece285.utils.evaluation import get_classification_accuracy

%matplotlib inline
plt.rcParams["figure.figsize"] = (10.0, 8.0) # set default size of plots

# For auto-reloading external modules
# See http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

# Use a subset of CIFAR10 for the assignment
dataset = get_cifar10_data(
    subset_train=5000,
    subset_val=250,
    subset_test=500,
)

print(dataset.keys())
print("Training Set Data Shape: ", dataset["x_train"].shape)
```

```

print("Training Set Label Shape: ", dataset["y_train"].shape)
print("Validation Set Data Shape: ", dataset["x_val"].shape)
print("Validation Set Label Shape: ", dataset["y_val"].shape)
print("Test Set Data Shape: ", dataset["x_test"].shape)
print("Test Set Label Shape: ", dataset["y_test"].shape)

```

```

dict_keys(['x_train', 'y_train', 'x_val', 'y_val', 'x_test', 'y_test'])
Training Set Data Shape: (5000, 3072)
Training Set Label Shape: (5000,)
Validation Set Data Shape: (250, 3072)
Validation Set Label Shape: (250,)
Test Set Data Shape: (500, 3072)
Test Set Label Shape: (500,)

```

```

[2]: x_train = dataset["x_train"]
      y_train = dataset["y_train"]
      x_val = dataset["x_val"]
      y_val = dataset["y_val"]
      x_test = dataset["x_test"]
      y_test = dataset["y_test"]

```

```

[3]: # Import more utilities and the layers you have implemented
      from ece285.layers.sequential import Sequential
      from ece285.layers.linear import Linear
      from ece285.layers.relu import ReLU
      from ece285.layers.softmax import Softmax
      from ece285.layers.loss_func import CrossEntropyLoss
      from ece285.utils.optimizer import SGD
      from ece285.utils.dataset import DataLoader
      from ece285.utils.trainer import Trainer

```

## 1.1 Visualize some examples from the dataset.

```

[4]: # We show a few examples of training images from each class.
      classes = [
          "airplane",
          "automobile",
          "bird",
          "cat",
          "deer",
          "dog",
          "frog",
          "horse",
          "ship",
      ]
      samples_per_class = 7

```

```

def visualize_data(dataset, classes, samples_per_class):
    num_classes = len(classes)
    for y, cls in enumerate(classes):
        idxs = np.flatnonzero(y_train == y)
        idxs = np.random.choice(idxs, samples_per_class, replace=False)
        for i, idx in enumerate(idxs):
            plt_idx = i * num_classes + y + 1
            plt.subplot(samples_per_class, num_classes, plt_idx)
            plt.imshow(dataset[idx])
            plt.axis("off")
            if i == 0:
                plt.title(cls)
    plt.show()

# Visualize the first 10 classes
visualize_data(
    x_train.reshape(5000, 3, 32, 32).transpose(0, 2, 3, 1),
    classes,
    samples_per_class,
)

```



## 1.2 Initialize the model

```
[5]: input_size = 3072
hidden_size = 100 # Hidden layer size (Hyper-parameter)
num_classes = 10 # Output

# For a default setting we use the same model we used for the toy dataset.
# This tells you the power of a 2 layered Neural Network. Recall the Universal
    ↳ Approximation Theorem.
# A 2 layer neural network with non-linearities can approximate any function,
    ↳ given large enough hidden layer
def init_model():
    # np.random.seed(0) # No need to fix the seed here
    l1 = Linear(input_size, hidden_size)
    l2 = Linear(hidden_size, num_classes)

    r1 = ReLU()
    softmax = Softmax()
```

```

return Sequential([l1, r1, l2, softmax])

# Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
net = init_model()
optim = SGD(net, lr=0.01, weight_decay=0.01)
loss_func = CrossEntropyLoss()
epoch = 200 # (Hyper-parameter)
batch_size = 200 # (Reduce the batch size if your computer is unable to handle
↳ it)

# Initialize the trainer class by passing the above modules
trainer = Trainer(
    dataset, optim, net, loss_func, epoch, batch_size, validate_interval=3
)

# Call the trainer function we have already implemented for you. This trains
↳ the model for the given
# hyper-parameters. It follows the same procedure as in the last ipython
↳ notebook you used for the toy-dataset
train_error, validation_accuracy = trainer.train()

```

```

Epoch Average Loss: 2.302537
Validate Acc: 0.084
Epoch Average Loss: 2.302358
Epoch Average Loss: 2.302153
Epoch Average Loss: 2.301861
Validate Acc: 0.104
Epoch Average Loss: 2.301433
Epoch Average Loss: 2.300851
Epoch Average Loss: 2.299964
Validate Acc: 0.096
Epoch Average Loss: 2.298798
Epoch Average Loss: 2.297321
Epoch Average Loss: 2.295501
Validate Acc: 0.084
Epoch Average Loss: 2.293346
Epoch Average Loss: 2.290893
Epoch Average Loss: 2.287849
Validate Acc: 0.084
Epoch Average Loss: 2.283949
Epoch Average Loss: 2.278915
Epoch Average Loss: 2.272740
Validate Acc: 0.096
Epoch Average Loss: 2.265747
Epoch Average Loss: 2.258297
Epoch Average Loss: 2.250794

```

Validate Acc: 0.100  
Epoch Average Loss: 2.243016  
Epoch Average Loss: 2.235507  
Epoch Average Loss: 2.228456  
Validate Acc: 0.116  
Epoch Average Loss: 2.221825  
Epoch Average Loss: 2.215697  
Epoch Average Loss: 2.210149  
Validate Acc: 0.124  
Epoch Average Loss: 2.204867  
Epoch Average Loss: 2.200080  
Epoch Average Loss: 2.195285  
Validate Acc: 0.128  
Epoch Average Loss: 2.191232  
Epoch Average Loss: 2.187177  
Epoch Average Loss: 2.183314  
Validate Acc: 0.136  
Epoch Average Loss: 2.179658  
Epoch Average Loss: 2.176573  
Epoch Average Loss: 2.173230  
Validate Acc: 0.144  
Epoch Average Loss: 2.170173  
Epoch Average Loss: 2.167557  
Epoch Average Loss: 2.164493  
Validate Acc: 0.148  
Epoch Average Loss: 2.161955  
Epoch Average Loss: 2.159274  
Epoch Average Loss: 2.156913  
Validate Acc: 0.144  
Epoch Average Loss: 2.154436  
Epoch Average Loss: 2.152519  
Epoch Average Loss: 2.150320  
Validate Acc: 0.144  
Epoch Average Loss: 2.148442  
Epoch Average Loss: 2.146027  
Epoch Average Loss: 2.144040  
Validate Acc: 0.152  
Epoch Average Loss: 2.142573  
Epoch Average Loss: 2.140663  
Epoch Average Loss: 2.138572  
Validate Acc: 0.152  
Epoch Average Loss: 2.137174  
Epoch Average Loss: 2.135517  
Epoch Average Loss: 2.133763  
Validate Acc: 0.148  
Epoch Average Loss: 2.132108  
Epoch Average Loss: 2.130843  
Epoch Average Loss: 2.128342

Validate Acc: 0.148  
Epoch Average Loss: 2.127817  
Epoch Average Loss: 2.125841  
Epoch Average Loss: 2.124409  
Validate Acc: 0.148  
Epoch Average Loss: 2.123041  
Epoch Average Loss: 2.121470  
Epoch Average Loss: 2.119940  
Validate Acc: 0.152  
Epoch Average Loss: 2.118505  
Epoch Average Loss: 2.116697  
Epoch Average Loss: 2.115945  
Validate Acc: 0.168  
Epoch Average Loss: 2.113718  
Epoch Average Loss: 2.112385  
Epoch Average Loss: 2.110330  
Validate Acc: 0.176  
Epoch Average Loss: 2.108615  
Epoch Average Loss: 2.106778  
Epoch Average Loss: 2.104724  
Validate Acc: 0.172  
Epoch Average Loss: 2.102461  
Epoch Average Loss: 2.100388  
Epoch Average Loss: 2.098450  
Validate Acc: 0.172  
Epoch Average Loss: 2.096084  
Epoch Average Loss: 2.093232  
Epoch Average Loss: 2.091584  
Validate Acc: 0.184  
Epoch Average Loss: 2.088124  
Epoch Average Loss: 2.085273  
Epoch Average Loss: 2.082583  
Validate Acc: 0.232  
Epoch Average Loss: 2.079115  
Epoch Average Loss: 2.076047  
Epoch Average Loss: 2.072967  
Validate Acc: 0.232  
Epoch Average Loss: 2.070173  
Epoch Average Loss: 2.067068  
Epoch Average Loss: 2.063438  
Validate Acc: 0.224  
Epoch Average Loss: 2.060540  
Epoch Average Loss: 2.057288  
Epoch Average Loss: 2.054034  
Validate Acc: 0.236  
Epoch Average Loss: 2.051270  
Epoch Average Loss: 2.048177  
Epoch Average Loss: 2.045248

Validate Acc: 0.240  
Epoch Average Loss: 2.042197  
Epoch Average Loss: 2.039321  
Epoch Average Loss: 2.036856  
Validate Acc: 0.240  
Epoch Average Loss: 2.034228  
Epoch Average Loss: 2.031098  
Epoch Average Loss: 2.029689  
Validate Acc: 0.260  
Epoch Average Loss: 2.026804  
Epoch Average Loss: 2.024594  
Epoch Average Loss: 2.021690  
Validate Acc: 0.244  
Epoch Average Loss: 2.019476  
Epoch Average Loss: 2.017448  
Epoch Average Loss: 2.015358  
Validate Acc: 0.260  
Epoch Average Loss: 2.013005  
Epoch Average Loss: 2.011135  
Epoch Average Loss: 2.008958  
Validate Acc: 0.268  
Epoch Average Loss: 2.007098  
Epoch Average Loss: 2.005913  
Epoch Average Loss: 2.003597  
Validate Acc: 0.268  
Epoch Average Loss: 2.001598  
Epoch Average Loss: 1.999838  
Epoch Average Loss: 1.997474  
Validate Acc: 0.268  
Epoch Average Loss: 1.995711  
Epoch Average Loss: 1.993893  
Epoch Average Loss: 1.992380  
Validate Acc: 0.272  
Epoch Average Loss: 1.990470  
Epoch Average Loss: 1.989180  
Epoch Average Loss: 1.987153  
Validate Acc: 0.276  
Epoch Average Loss: 1.985827  
Epoch Average Loss: 1.983262  
Epoch Average Loss: 1.981728  
Validate Acc: 0.280  
Epoch Average Loss: 1.980198  
Epoch Average Loss: 1.978801  
Epoch Average Loss: 1.976564  
Validate Acc: 0.288  
Epoch Average Loss: 1.974345  
Epoch Average Loss: 1.972096  
Epoch Average Loss: 1.970484



Validate Acc: 0.288  
Epoch Average Loss: 1.967622  
Epoch Average Loss: 1.965336  
Epoch Average Loss: 1.963110  
Validate Acc: 0.304  
Epoch Average Loss: 1.961092  
Epoch Average Loss: 1.958728  
Epoch Average Loss: 1.954832  
Validate Acc: 0.308  
Epoch Average Loss: 1.951695  
Epoch Average Loss: 1.949919  
Epoch Average Loss: 1.946959  
Validate Acc: 0.300  
Epoch Average Loss: 1.943884  
Epoch Average Loss: 1.940049  
Epoch Average Loss: 1.937556  
Validate Acc: 0.288  
Epoch Average Loss: 1.933858  
Epoch Average Loss: 1.931535  
Epoch Average Loss: 1.928331  
Validate Acc: 0.292  
Epoch Average Loss: 1.925393  
Epoch Average Loss: 1.923706  
Epoch Average Loss: 1.920337  
Validate Acc: 0.292  
Epoch Average Loss: 1.918130  
Epoch Average Loss: 1.914838  
Epoch Average Loss: 1.914686  
Validate Acc: 0.296  
Epoch Average Loss: 1.911120  
Epoch Average Loss: 1.910004  
Epoch Average Loss: 1.907682  
Validate Acc: 0.284  
Epoch Average Loss: 1.903876  
Epoch Average Loss: 1.902066  
Epoch Average Loss: 1.899852  
Validate Acc: 0.292  
Epoch Average Loss: 1.898004  
Epoch Average Loss: 1.893862  
Epoch Average Loss: 1.894274  
Validate Acc: 0.312  
Epoch Average Loss: 1.891401  
Epoch Average Loss: 1.888412  
Epoch Average Loss: 1.887628  
Validate Acc: 0.304  
Epoch Average Loss: 1.883805  
Epoch Average Loss: 1.883570  
Epoch Average Loss: 1.883284

Validate Acc: 0.308  
Epoch Average Loss: 1.879298  
Epoch Average Loss: 1.876486  
Epoch Average Loss: 1.876580  
Validate Acc: 0.304  
Epoch Average Loss: 1.873125  
Epoch Average Loss: 1.872160  
Epoch Average Loss: 1.868687  
Validate Acc: 0.292  
Epoch Average Loss: 1.868044  
Epoch Average Loss: 1.866335  
Epoch Average Loss: 1.863730  
Validate Acc: 0.284  
Epoch Average Loss: 1.862398  
Epoch Average Loss: 1.859831  
Epoch Average Loss: 1.857877  
Validate Acc: 0.292  
Epoch Average Loss: 1.856972  
Epoch Average Loss: 1.854295  
Epoch Average Loss: 1.852176  
Validate Acc: 0.300  
Epoch Average Loss: 1.850807  
Epoch Average Loss: 1.848620  
Epoch Average Loss: 1.846455  
Validate Acc: 0.296  
Epoch Average Loss: 1.844929  
Epoch Average Loss: 1.843053  
Epoch Average Loss: 1.838885  
Validate Acc: 0.296  
Epoch Average Loss: 1.838025  
Epoch Average Loss: 1.835789  
Epoch Average Loss: 1.834729  
Validate Acc: 0.324  
Epoch Average Loss: 1.832280  
Epoch Average Loss: 1.830350  
Epoch Average Loss: 1.828094  
Validate Acc: 0.304  
Epoch Average Loss: 1.826341  
Epoch Average Loss: 1.822406  
Epoch Average Loss: 1.821506  
Validate Acc: 0.312  
Epoch Average Loss: 1.820259  
Epoch Average Loss: 1.816891  
Epoch Average Loss: 1.816177  
Validate Acc: 0.320  
Epoch Average Loss: 1.815976  
Epoch Average Loss: 1.811736  
Epoch Average Loss: 1.811231

Validate Acc: 0.324  
Epoch Average Loss: 1.811461

### 1.2.1 Print the training and validation accuracies for the default hyper-parameters provided

```
[6]: from ece285.utils.evaluation import get_classification_accuracy

out_train = net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)
```

Training acc: 0.3426  
Validation acc: 0.328

### 1.2.2 Debug the training

With the default parameters we provided above, you should get a validation accuracy of around ~0.2 on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the training loss function and the validation accuracies during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[7]: # Plot the training loss function and validation accuracies
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()
```



```
[8]: from ece285.utils.vis_utils import visualize_grid

# Credits: http://cs231n.stanford.edu/

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype("uint8"))
    plt.gca().axis("off")
    plt.show()

show_net_weights(net)
```



## 2 Tune your hyperparameters (50%)

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 40% on the validation set. Our best network gets over 40% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on cifar10 as you can (40% could serve as a reference), with a fully-connected Neural Network.

**Explain your hyperparameter tuning process below.**

**Your Answer:** Increase the learning rate to 0.2, and find the training loss does not decrease, so I search and tried the learning rate between 0.02 and 0.1.

I also tried to increase the hidden layer size to 512, however, it seems to be overfitting.

I also noticed that there is a gap between the training and validation accuracy, so I tried to add weight decay to 0.015.

Finally, I noticed the training loss is still decreasing for 200 epoch, so I increased the epoch to 400.

```
[16]: input_size = 3072
hidden_size = 300 # Hidden layer size (Hyper-parameter)
num_classes = 10 # Output

# For a default setting we use the same model we used for the toy dataset.
# This tells you the power of a 2 layered Neural Network. Recall the Universal
↳ Approximation Theorem.
# A 2 layer neural network with non-linearities can approximate any function,
↳ given large enough hidden layer
def tuned_model():
    # np.random.seed(0) # No need to fix the seed here
    l1 = Linear(input_size, hidden_size)
    # l2 = Linear(hidden_size, hidden_size)
    l3 = Linear(hidden_size, num_classes)

    r1 = ReLU()
    # r2 = ReLU()
    softmax = Softmax()
    return Sequential([l1, r1, l3, softmax])

# Initialize the dataset with the dataloader class
dataset = DataLoader(x_train, y_train, x_val, y_val, x_test, y_test)
best_net = tuned_model()
optim = SGD(best_net, lr=0.065, weight_decay=0.015)
loss_func = CrossEntropyLoss()
epoch = 350 # (Hyper-parameter)
batch_size = 256 # (Reduce the batch size if your computer is unable to handle
↳ it)

# Initialize the trainer class by passing the above modules
trainer = Trainer(
    dataset, optim, best_net, loss_func, epoch, batch_size, validate_interval=5
)
```

```
# Call the trainer function we have already implemented for you. This trains  
↳ the model for the given  
# hyper-parameters. It follows the same procedure as in the last ipython  
↳ notebook you used for the toy-dataset  
train_error, validation_accuracy = trainer.train()
```

```
Epoch Average Loss: 2.301753  
Validate Acc: 0.104  
Epoch Average Loss: 2.296516  
Epoch Average Loss: 2.282466  
Epoch Average Loss: 2.248706  
Epoch Average Loss: 2.216719  
Epoch Average Loss: 2.190613  
Validate Acc: 0.140  
Epoch Average Loss: 2.178611  
Epoch Average Loss: 2.166060  
Epoch Average Loss: 2.151628  
Epoch Average Loss: 2.143922  
Epoch Average Loss: 2.139498  
Validate Acc: 0.168  
Epoch Average Loss: 2.127751  
Epoch Average Loss: 2.126133  
Epoch Average Loss: 2.103754  
Epoch Average Loss: 2.103363  
Epoch Average Loss: 2.081593  
Validate Acc: 0.176  
Epoch Average Loss: 2.071344  
Epoch Average Loss: 2.046820  
Epoch Average Loss: 2.043875  
Epoch Average Loss: 2.027386  
Epoch Average Loss: 2.040150  
Validate Acc: 0.276  
Epoch Average Loss: 2.037707  
Epoch Average Loss: 2.015576  
Epoch Average Loss: 2.021603  
Epoch Average Loss: 2.001865  
Epoch Average Loss: 1.983514  
Validate Acc: 0.276  
Epoch Average Loss: 1.999110  
Epoch Average Loss: 2.007294  
Epoch Average Loss: 1.956605  
Epoch Average Loss: 1.940943  
Epoch Average Loss: 1.938132  
Validate Acc: 0.284  
Epoch Average Loss: 1.923839  
Epoch Average Loss: 1.918498
```

Epoch Average Loss: 1.926222  
Epoch Average Loss: 1.932101  
Epoch Average Loss: 1.908986  
Validate Acc: 0.260  
Epoch Average Loss: 1.903900  
Epoch Average Loss: 1.905333  
Epoch Average Loss: 1.871296  
Epoch Average Loss: 1.878767  
Epoch Average Loss: 1.868555  
Validate Acc: 0.292  
Epoch Average Loss: 1.876370  
Epoch Average Loss: 1.844241  
Epoch Average Loss: 1.879067  
Epoch Average Loss: 1.853410  
Epoch Average Loss: 1.814960  
Validate Acc: 0.328  
Epoch Average Loss: 1.838543  
Epoch Average Loss: 1.830557  
Epoch Average Loss: 1.814447  
Epoch Average Loss: 1.799894  
Epoch Average Loss: 1.837435  
Validate Acc: 0.316  
Epoch Average Loss: 1.806856  
Epoch Average Loss: 1.801321  
Epoch Average Loss: 1.776798  
Epoch Average Loss: 1.782245  
Epoch Average Loss: 1.780860  
Validate Acc: 0.308  
Epoch Average Loss: 1.779238  
Epoch Average Loss: 1.782210  
Epoch Average Loss: 1.782986  
Epoch Average Loss: 1.768374  
Epoch Average Loss: 1.764355  
Validate Acc: 0.368  
Epoch Average Loss: 1.761973  
Epoch Average Loss: 1.758296  
Epoch Average Loss: 1.753897  
Epoch Average Loss: 1.767010  
Epoch Average Loss: 1.744071  
Validate Acc: 0.368  
Epoch Average Loss: 1.729908  
Epoch Average Loss: 1.741825  
Epoch Average Loss: 1.723569  
Epoch Average Loss: 1.755336  
Epoch Average Loss: 1.741392  
Validate Acc: 0.320  
Epoch Average Loss: 1.707641  
Epoch Average Loss: 1.746871



Epoch Average Loss: 1.722200  
Epoch Average Loss: 1.732470  
Epoch Average Loss: 1.710167  
Validate Acc: 0.316  
Epoch Average Loss: 1.729676  
Epoch Average Loss: 1.731321  
Epoch Average Loss: 1.709032  
Epoch Average Loss: 1.723418  
Epoch Average Loss: 1.724400  
Validate Acc: 0.368  
Epoch Average Loss: 1.695499  
Epoch Average Loss: 1.695465  
Epoch Average Loss: 1.719601  
Epoch Average Loss: 1.713343  
Epoch Average Loss: 1.685659  
Validate Acc: 0.316  
Epoch Average Loss: 1.718871  
Epoch Average Loss: 1.667153  
Epoch Average Loss: 1.716438  
Epoch Average Loss: 1.680555  
Epoch Average Loss: 1.687059  
Validate Acc: 0.364  
Epoch Average Loss: 1.684935  
Epoch Average Loss: 1.701476  
Epoch Average Loss: 1.669628  
Epoch Average Loss: 1.689408  
Epoch Average Loss: 1.668413  
Validate Acc: 0.360  
Epoch Average Loss: 1.670627  
Epoch Average Loss: 1.669474  
Epoch Average Loss: 1.648149  
Epoch Average Loss: 1.668217  
Epoch Average Loss: 1.661556  
Validate Acc: 0.328  
Epoch Average Loss: 1.660184  
Epoch Average Loss: 1.672420  
Epoch Average Loss: 1.668235  
Epoch Average Loss: 1.624360  
Epoch Average Loss: 1.665213  
Validate Acc: 0.396  
Epoch Average Loss: 1.657199  
Epoch Average Loss: 1.684161  
Epoch Average Loss: 1.652947  
Epoch Average Loss: 1.663393  
Epoch Average Loss: 1.641019  
Validate Acc: 0.388  
Epoch Average Loss: 1.618663  
Epoch Average Loss: 1.655396

Epoch Average Loss: 1.630275  
Epoch Average Loss: 1.647604  
Epoch Average Loss: 1.651459  
Validate Acc: 0.344  
Epoch Average Loss: 1.655329  
Epoch Average Loss: 1.613166  
Epoch Average Loss: 1.624466  
Epoch Average Loss: 1.636580  
Epoch Average Loss: 1.629607  
Validate Acc: 0.296  
Epoch Average Loss: 1.641421  
Epoch Average Loss: 1.607866  
Epoch Average Loss: 1.634251  
Epoch Average Loss: 1.607163  
Epoch Average Loss: 1.655783  
Validate Acc: 0.408  
Epoch Average Loss: 1.608675  
Epoch Average Loss: 1.613215  
Epoch Average Loss: 1.622214  
Epoch Average Loss: 1.577627  
Epoch Average Loss: 1.594959  
Validate Acc: 0.312  
Epoch Average Loss: 1.622612  
Epoch Average Loss: 1.609518  
Epoch Average Loss: 1.623246  
Epoch Average Loss: 1.609658  
Epoch Average Loss: 1.618266  
Validate Acc: 0.400  
Epoch Average Loss: 1.609498  
Epoch Average Loss: 1.570585  
Epoch Average Loss: 1.569558  
Epoch Average Loss: 1.613508  
Epoch Average Loss: 1.614992  
Validate Acc: 0.380  
Epoch Average Loss: 1.609732  
Epoch Average Loss: 1.576295  
Epoch Average Loss: 1.606699  
Epoch Average Loss: 1.580100  
Epoch Average Loss: 1.594790  
Validate Acc: 0.364  
Epoch Average Loss: 1.582620  
Epoch Average Loss: 1.618606  
Epoch Average Loss: 1.591107  
Epoch Average Loss: 1.616387  
Epoch Average Loss: 1.612133  
Validate Acc: 0.368  
Epoch Average Loss: 1.561637  
Epoch Average Loss: 1.587011

Epoch Average Loss: 1.539004  
Epoch Average Loss: 1.592941  
Epoch Average Loss: 1.582667  
Validate Acc: 0.360  
Epoch Average Loss: 1.589324  
Epoch Average Loss: 1.545207  
Epoch Average Loss: 1.567754  
Epoch Average Loss: 1.585496  
Epoch Average Loss: 1.562687  
Validate Acc: 0.376  
Epoch Average Loss: 1.579881  
Epoch Average Loss: 1.545406  
Epoch Average Loss: 1.584983  
Epoch Average Loss: 1.597885  
Epoch Average Loss: 1.589558  
Validate Acc: 0.368  
Epoch Average Loss: 1.565554  
Epoch Average Loss: 1.574760  
Epoch Average Loss: 1.570246  
Epoch Average Loss: 1.536783  
Epoch Average Loss: 1.566466  
Validate Acc: 0.404  
Epoch Average Loss: 1.568723  
Epoch Average Loss: 1.558358  
Epoch Average Loss: 1.562357  
Epoch Average Loss: 1.557687  
Epoch Average Loss: 1.543685  
Validate Acc: 0.408  
Epoch Average Loss: 1.571142  
Epoch Average Loss: 1.583448  
Epoch Average Loss: 1.590934  
Epoch Average Loss: 1.580021  
Epoch Average Loss: 1.543568  
Validate Acc: 0.424  
Epoch Average Loss: 1.550869  
Epoch Average Loss: 1.549791  
Epoch Average Loss: 1.578133  
Epoch Average Loss: 1.583584  
Epoch Average Loss: 1.534206  
Validate Acc: 0.416  
Epoch Average Loss: 1.542186  
Epoch Average Loss: 1.587810  
Epoch Average Loss: 1.551319  
Epoch Average Loss: 1.529754  
Epoch Average Loss: 1.523737  
Validate Acc: 0.380  
Epoch Average Loss: 1.577863  
Epoch Average Loss: 1.543801

Epoch Average Loss: 1.546237  
Epoch Average Loss: 1.494706  
Epoch Average Loss: 1.541451  
Validate Acc: 0.368  
Epoch Average Loss: 1.536401  
Epoch Average Loss: 1.522190  
Epoch Average Loss: 1.548335  
Epoch Average Loss: 1.548021  
Epoch Average Loss: 1.524689  
Validate Acc: 0.428  
Epoch Average Loss: 1.543772  
Epoch Average Loss: 1.528681  
Epoch Average Loss: 1.531962  
Epoch Average Loss: 1.519973  
Epoch Average Loss: 1.523594  
Validate Acc: 0.356  
Epoch Average Loss: 1.560731  
Epoch Average Loss: 1.525214  
Epoch Average Loss: 1.531816  
Epoch Average Loss: 1.523027  
Epoch Average Loss: 1.561042  
Validate Acc: 0.368  
Epoch Average Loss: 1.531487  
Epoch Average Loss: 1.542539  
Epoch Average Loss: 1.530839  
Epoch Average Loss: 1.546576  
Epoch Average Loss: 1.523077  
Validate Acc: 0.360  
Epoch Average Loss: 1.566358  
Epoch Average Loss: 1.531781  
Epoch Average Loss: 1.514562  
Epoch Average Loss: 1.510652  
Epoch Average Loss: 1.537120  
Validate Acc: 0.376  
Epoch Average Loss: 1.564607  
Epoch Average Loss: 1.504075  
Epoch Average Loss: 1.547519  
Epoch Average Loss: 1.522866  
Epoch Average Loss: 1.499864  
Validate Acc: 0.424  
Epoch Average Loss: 1.544818  
Epoch Average Loss: 1.521793  
Epoch Average Loss: 1.518819  
Epoch Average Loss: 1.526504  
Epoch Average Loss: 1.501296  
Validate Acc: 0.412  
Epoch Average Loss: 1.513195  
Epoch Average Loss: 1.530199

Epoch Average Loss: 1.500367  
Epoch Average Loss: 1.509301  
Epoch Average Loss: 1.498574  
Validate Acc: 0.416  
Epoch Average Loss: 1.511133  
Epoch Average Loss: 1.502791  
Epoch Average Loss: 1.532824  
Epoch Average Loss: 1.477778  
Epoch Average Loss: 1.499101  
Validate Acc: 0.424  
Epoch Average Loss: 1.462172  
Epoch Average Loss: 1.528518  
Epoch Average Loss: 1.551203  
Epoch Average Loss: 1.513319  
Epoch Average Loss: 1.501783  
Validate Acc: 0.428  
Epoch Average Loss: 1.525013  
Epoch Average Loss: 1.500367  
Epoch Average Loss: 1.521106  
Epoch Average Loss: 1.472512  
Epoch Average Loss: 1.512831  
Validate Acc: 0.392  
Epoch Average Loss: 1.518030  
Epoch Average Loss: 1.495580  
Epoch Average Loss: 1.464169  
Epoch Average Loss: 1.492667  
Epoch Average Loss: 1.495597  
Validate Acc: 0.404  
Epoch Average Loss: 1.474860  
Epoch Average Loss: 1.497158  
Epoch Average Loss: 1.507914  
Epoch Average Loss: 1.503335  
Epoch Average Loss: 1.473537  
Validate Acc: 0.348  
Epoch Average Loss: 1.495819  
Epoch Average Loss: 1.499314  
Epoch Average Loss: 1.468144  
Epoch Average Loss: 1.485549  
Epoch Average Loss: 1.499831  
Validate Acc: 0.384  
Epoch Average Loss: 1.485386  
Epoch Average Loss: 1.532145  
Epoch Average Loss: 1.498870  
Epoch Average Loss: 1.498958  
Epoch Average Loss: 1.483886  
Validate Acc: 0.420  
Epoch Average Loss: 1.483339  
Epoch Average Loss: 1.502875

Epoch Average Loss: 1.480103  
Epoch Average Loss: 1.523304  
Epoch Average Loss: 1.564988  
Validate Acc: 0.396  
Epoch Average Loss: 1.497646  
Epoch Average Loss: 1.512994  
Epoch Average Loss: 1.492358  
Epoch Average Loss: 1.509864  
Epoch Average Loss: 1.503724  
Validate Acc: 0.384  
Epoch Average Loss: 1.483003  
Epoch Average Loss: 1.506733  
Epoch Average Loss: 1.452618  
Epoch Average Loss: 1.503709  
Epoch Average Loss: 1.492359  
Validate Acc: 0.324  
Epoch Average Loss: 1.511281  
Epoch Average Loss: 1.492560  
Epoch Average Loss: 1.506510  
Epoch Average Loss: 1.494733  
Epoch Average Loss: 1.472759  
Validate Acc: 0.408  
Epoch Average Loss: 1.486197  
Epoch Average Loss: 1.477468  
Epoch Average Loss: 1.487132  
Epoch Average Loss: 1.477691  
Epoch Average Loss: 1.460404  
Validate Acc: 0.428  
Epoch Average Loss: 1.472240  
Epoch Average Loss: 1.492068  
Epoch Average Loss: 1.481140  
Epoch Average Loss: 1.452341  
Epoch Average Loss: 1.498336  
Validate Acc: 0.424  
Epoch Average Loss: 1.500220  
Epoch Average Loss: 1.473680  
Epoch Average Loss: 1.456573  
Epoch Average Loss: 1.481263  
Epoch Average Loss: 1.463041  
Validate Acc: 0.412  
Epoch Average Loss: 1.473260  
Epoch Average Loss: 1.484846  
Epoch Average Loss: 1.482217  
Epoch Average Loss: 1.463598  
Epoch Average Loss: 1.474441  
Validate Acc: 0.380  
Epoch Average Loss: 1.470875  
Epoch Average Loss: 1.464585

Epoch Average Loss: 1.479795  
Epoch Average Loss: 1.484590  
Epoch Average Loss: 1.497924  
Validate Acc: 0.408  
Epoch Average Loss: 1.456147  
Epoch Average Loss: 1.471779  
Epoch Average Loss: 1.455268  
Epoch Average Loss: 1.469981  
Epoch Average Loss: 1.485015  
Validate Acc: 0.372  
Epoch Average Loss: 1.448128  
Epoch Average Loss: 1.478286  
Epoch Average Loss: 1.458857  
Epoch Average Loss: 1.465127  
Epoch Average Loss: 1.457996  
Validate Acc: 0.392  
Epoch Average Loss: 1.481294  
Epoch Average Loss: 1.466897  
Epoch Average Loss: 1.453162  
Epoch Average Loss: 1.471954  
Epoch Average Loss: 1.452563  
Validate Acc: 0.420  
Epoch Average Loss: 1.462648  
Epoch Average Loss: 1.468751  
Epoch Average Loss: 1.466728  
Epoch Average Loss: 1.464488  
Epoch Average Loss: 1.464801  
Validate Acc: 0.416  
Epoch Average Loss: 1.460953  
Epoch Average Loss: 1.495575  
Epoch Average Loss: 1.440437  
Epoch Average Loss: 1.460605  
Epoch Average Loss: 1.482682  
Validate Acc: 0.356  
Epoch Average Loss: 1.440953  
Epoch Average Loss: 1.504439  
Epoch Average Loss: 1.520626  
Epoch Average Loss: 1.461737  
Epoch Average Loss: 1.465475  
Validate Acc: 0.432  
Epoch Average Loss: 1.465432  
Epoch Average Loss: 1.445655  
Epoch Average Loss: 1.467228  
Epoch Average Loss: 1.483421

```
[17]: from ece285.utils.evaluation import get_classification_accuracy
```


```
out_train = best_net.predict(x_train)
acc = get_classification_accuracy(out_train, y_train)
print("Training acc: ", acc)
out_val = best_net.predict(x_val)
acc = get_classification_accuracy(out_val, y_val)
print("Validation acc: ", acc)
```

Training acc: 0.478  
Validation acc: 0.412

```
[18]: # Plot the training loss function and validation accuracies
plt.subplot(2, 1, 1)
plt.plot(train_error)
plt.title("Training Loss History")
plt.xlabel("Iteration")
plt.ylabel("Loss")

plt.subplot(2, 1, 2)
# plt.plot(stats['train_acc_history'], label='train')
plt.plot(validation_accuracy, label="val")
plt.title("Classification accuracy history")
plt.xlabel("Epoch")
plt.ylabel("Classification accuracy")
plt.legend()
plt.show()
```






```
[19]: from ece285.utils.vis_utils import visualize_grid

# Credits: http://cs231n.stanford.edu/

# Visualize the weights of the network

def show_net_weights(net):
    W1 = net._modules[0].parameters[0]
    W1 = W1.reshape(3, 32, 32, -1).transpose(3, 1, 2, 0)
    plt.imshow(visualize_grid(W1, padding=3).astype("uint8"))
    plt.gca().axis("off")
    plt.show()

show_net_weights(best_net)
```



```
[20]: best_net_hyperparams = [0.065, 0.015, 350, 300]
#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
#
# model hyperparams in best_net.
#
#
# To help debug your network, it may help to use visualizations similar to the
#
# ones we used above; these visualizations will have significant qualitative
```

```

# differences from the ones we saw above for the poorly tuned network.
↪#
#
↪#
# You are now free to test different combinations of hyperparameters to build
↪#
# various models and test them according to the above plots and visualization
↪#

# TODO: Show the above plots and visualizations for the default params (already
↪#
# done) and the best hyper-params you obtain. You only need to show this for 2
↪#
# sets of hyper-params.
↪#
# You just need to store values for the hyperparameters in best_net_hyperparams
↪#
# as a list in the order
# best_net_hyperparams = [lr, weight_decay, epoch, hidden_size]
#####
pass

```

```

[21]: # TODO: Plot the training_error and validation_accuracy of the best network (5%)

# TODO: visualize the weights of the best network (5%)

```

### 3 Run on the test set (30%)

When you are done experimenting, you should evaluate your final trained network on the test set; you should get above 35%.

```

[22]: test_acc = (best_net.predict(x_test) == y_test).mean()
print("Test accuracy: ", test_acc)

```

Test accuracy: 0.356

**Inline Question (10%)** Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Select all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

**Your Answer:** The correct options to help decrease the gap between training and testing accuracy in a neural network classifier are: ##### Your Explanation: Train on a larger dataset - Training on a larger dataset can help the model generalize better to unseen data by providing more examples from which it can learn. This reduces overfitting, which is a common reason for the discrepancy between training and testing performance. Increase the regularization strength - Regularization techniques such as L1 or L2 regularization, or using dropout, help prevent the model from fitting too closely to the noise in the training data. By increasing the regularization strength, the model is encouraged to develop simpler, more generalizable patterns, thus reducing overfitting. Option 2, adding more hidden units, is generally not recommended to decrease overfitting, as it can actually lead to more complex models that overfit the training data even more. Thus, it might increase the gap between training and testing accuracy rather than decrease it.

Therefore, the best options are 1 and 3.

[15] :